

Credit Card Fault Detection

Model Development and Deployment Process

Date	09-01-2024
Prepared by	Makineedi Venkat Dinesh

Fraud Prediction

Document details the process of estimating fault transaction based on machine learning approach.

Author	Makineedi Venkat Dinesh	Date – 09/01/2024
--------	-------------------------	-------------------

Table of Content

1) Prediction Approach	2
2) Model Deployment (AWS sagemaker & AWS API Gateway)	24
3) Future Roadmap	4

Credit Card Fault Detection Report

1) Prediction Approach

I experimented with three approaches and finalized one approach giving best area under ROC curve.

Approach 1: Normalizing 'Amount' column and eliminating 'time' column.

Approach 2: Normalizing all features.

Approach 3: Resampling all the data to get balanced dataset.

In every approach I trained three models:

Model1: lightGBM -> hyperparameter tuning is done for this model.

Model2: Logistic Regression

Model3: KNeighbours classification

Approach 1: Normalizing 'Amount' column and eliminating 'time' column

Model1: lightGBM reports:

TRAIN MODEL CLASSIFICATION REPORT





	precision	recall	f1-score	support
No Fraud	1.00	1.00	1.00	227451
Fraud	1.00	1.00	1.00	394
accuracy			1.00	227845

TEST MODEL CLASSIFICATION REPORT

	precision	recall	f1-score	support
No Fraud	1.00	1.00	1.00	56864
Fraud	0.98	0.81	0.88	98
accuracy			1.00	56962

Model performance report

lgbm_clf_tuned



Model2: Logistic Regression reports:

TRAIN MODEL CLASSIFICATION REPORT

	precision	recall	f1-score	support
No Fraud	1.00	1.00	1.00	227451
Fraud	0.87	0.62	0.72	394
accuracy			1.00	227845

TEST MODEL CLASSIFICATION REPORT

	precision	recall	f1-score	support
No Fraud	1.00	1.00	1.00	56864
Fraud	0.92	0.61	0.74	98
accuracy			1.00	56962


Model performance report

lr


Confusion Matrix

	0 (No Fraud) pred	1 (Fraud) pred
1 (Fraud) actual	38	60
0 (No Fraud) actual	56859	5


Metrics




ROC curve (0.98)



Precision - Recall curve



Cumulative gains curve



Model3: KNeighbors Classification reports:

TRAIN MODEL CLASSIFICATION REPORT





	precision	recall	f1-score	support
No Fraud	1.00	1.00	1.00	227451
Fraud	0.98	0.79	0.87	394
accuracy			1.00	227845

TEST MODEL CLASSIFICATION REPORT


	precision	recall	f1-score	support
No Fraud	1.00	1.00	1.00	56864
Fraud	0.97	0.78	0.86	98
accuracy			1.00	56962

Model performance report

knn



ROC Curve for lightgbm, Logistic Regression and K-Neighbors Models:



From the above reports and ROC curve, lightGBM model is having highest ROC area of 0.98.

Approach 2: Normalizing all features.

Model1: lightGBM reports:

TRAIN MODEL CLASSIFICATION REPORT





	precision	recall	f1-score	support
No Fraud	1.00	1.00	1.00	227468
Fraud	1.00	1.00	1.00	377
accuracy			1.00	227845

TEST MODEL CLASSIFICATION REPORT

	precision	recall	f1-score	support
No Fraud	1.00	1.00	1.00	56847
Fraud	0.93	0.83	0.88	115
accuracy			1.00	56962

Model performance report

lgbm_clf_tuned



Model2: Logistic Regression reports:

TRAIN MODEL CLASSIFICATION REPORT

	precision	recall	f1-score	support
No Fraud	1.00	1.00	1.00	227468
Fraud	0.89	0.63	0.74	377
accuracy			1.00	227845

TEST MODEL CLASSIFICATION REPORT

	precision	recall	f1-score	support
No Fraud	1.00	1.00	1.00	56847
Fraud	0.83	0.61	0.70	115
accuracy			1.00	56962

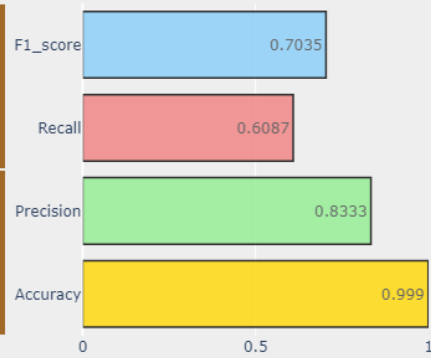
Model performance report

lr


Confusion Matrix

	0 (No Fraud) pred	1 (Fraud) pred
1 (Fraud) actual	45	70
0 (No Fraud) actual	56833	14


Metrics




ROC curve (0.977)



Precision - Recall curve



Cumulative gains curve



Model3: KNeighbors Classification reports:

TRAIN MODEL CLASSIFICATION REPORT


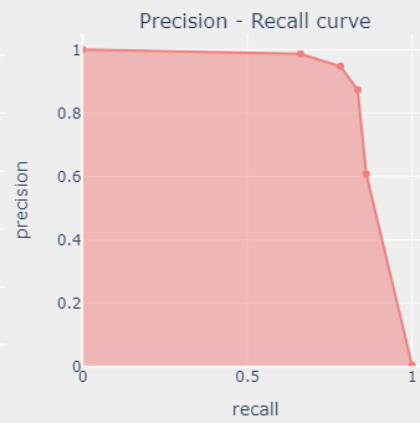


	precision	recall	f1-score	support
No Fraud	1.00	1.00	1.00	227468
Fraud	0.97	0.79	0.87	377
accuracy			1.00	227845

TEST MODEL CLASSIFICATION REPORT


	precision	recall	f1-score	support
No Fraud	1.00	1.00	1.00	56847
Fraud	0.95	0.78	0.86	115
accuracy			1.00	56962

Model performance report

knn



ROC Curve for lightgbm, Logistic Regression and K-Neighbors Models:



Approach 3: Resampling.

Model1: lightGBM reports:

TRAIN MODEL CLASSIFICATION REPORT





	precision	recall	f1-score	support
No Fraud	0.93	0.98	0.96	383
Fraud	0.98	0.93	0.96	404
accuracy			0.96	787

TEST MODEL CLASSIFICATION REPORT

	precision	recall	f1-score	support
No Fraud	0.93	0.94	0.94	109
Fraud	0.93	0.91	0.92	88
accuracy			0.93	197

Model performance report

lgbm_clf_tuned_b



Model2: Logistic Regression reports:

TRAIN MODEL CLASSIFICATION REPORT

	precision	recall	f1-score	support
No Fraud	0.92	0.98	0.95	383
Fraud	0.98	0.92	0.95	404
accuracy			0.95	787

TEST MODEL CLASSIFICATION REPORT

	precision	recall	f1-score	support
No Fraud	0.94	0.97	0.95	109
Fraud	0.96	0.92	0.94	88
accuracy			0.95	197

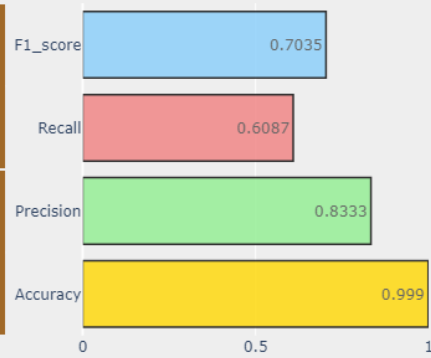
Model performance report

lr


Confusion Matrix

	0 (No Fraud) pred	1 (Fraud) pred
1 (Fraud) actual	45	70
0 (No Fraud) actual	56833	14


Metrics




ROC curve (0.977)



Precision - Recall curve



Cumulative gains curve



Model3: KNeighbors Classification reports:

TRAIN MODEL CLASSIFICATION REPORT





	precision	recall	f1-score	support
No Fraud	0.88	0.99	0.93	383
Fraud	0.99	0.87	0.93	404
accuracy			0.93	787

TEST MODEL CLASSIFICATION REPORT


	precision	recall	f1-score	support
No Fraud	0.89	1.00	0.94	109
Fraud	1.00	0.84	0.91	88
accuracy			0.93	197

Model performance report

knn_b



ROC Curve for lightgbm, Logistic Regression and K-Neighbors Models:



CONCLUSION:

Depending on the specific requirements of problem (e.g., the importance of false positives vs. false negatives), we may choose one model over the other based on the balance between precision and recall. Considering the context of Credit Card Fault Prediction application misclassifications are more costly than others. Hence I am selecting the Approach-3: Resampling and Normalizing all features and lightGBM model as this model is giving highest Recall of 0.9091 of among all others and good ROC area of 0.972.

So, I will be using that model to **deploy** it in **AWS sagemaker** and integrate the endpoint with **AWS API Gateway** to publicly access the Invoke endpoint.

2)Model Deployment Process

Docker Image Creation Process for the developed model:

- 1) Create a folder named "service" containing 3 python scripts (predictor.py, service.py, serve) and nginx.conf file for model deployment
- 2) Build a docker image of service folder with 'Dockerfile.sagemaker' using the following command:
docker build . -f Dockerfile.sagemaker -t <image_name>:<tag_name>

NOTE: In current deployment, <image_name> = faultpredictor and <tag_name> = im

- 3) Add AWS ECR credentials using following command: aws ecr get-login-password --region us-east-1 | docker login --username AWS --password-stdin <account_id>.dkr.ecr.us-east-1.amazonaws.com
- 4) If in case docker image_name and tag_name has to be changed use the following command:
docker tag <previous_image_name>:<previous_tag_name> <account_id>.dkr.ecr.us-east-1.amazonaws.com/<latest_image_name>:<latest_tag_name>
- 5) Push the image to AWS ECR using the following command:
docker push <account_id>.dkr.ecr.us-east-1.amazonaws.com/<latest_image_name>:<latest_tag_name>

NOTE: In current deployment, <latest_image_name> = faultpredictor and <latest_tag_name> = cc

After Docker Image is pushed to ECR run a "deploy.py":

- 6) Now deploy the model in sagemaker and create an endpoint in sagemaker to invoke the model using deploy.py file using the following command:
python deploy.py --model_data <S3 URI of model location> --image_uri <AWS ECR image uri> --model_server_timeout <timeout> --endpoint <endpoint_name>

NOTE: In current deployment:

<S3 URI of model location> = "s3://keysec/cc-experiments/local-developed-model/model/im/model.tar.gz"

<AWS ECR image uri> = "<account_id>.dkr.ecr.us-east-1.amazonaws.com/faultpredictor:im"

<timeout> = 300

<endpoint_name> = "im"

After running "deploy.py" file, Integration of created sagemaker endpoint with API Gateway:

7)

a) In AWS API Gateway select "sagemaker" API

<input type="radio"/>	sagemaker	calling sagemaker endpoint	fluqdn751l	REST	Regional	2023-11-21
-----------------------	---------------------------	----------------------------	------------	------	----------	------------

b) In "sagemaker" API, create resource with resource path and resource name

[API Gateway](#) > [APIs](#) > [Resources - sagemaker \(fluqdn751l\)](#) > [Create resource](#)

Create resource

Resource details

☒ Proxy resource [Info](#)
Proxy resources handle requests to all sub-resources. To create a proxy resource use a path parameter that ends with a plus sign, for example {proxy+}.

Resource path:

Resource name:

☐ CORS (Cross Origin Resource Sharing) [Info](#)
Create an OPTIONS method that allows all origins, all methods, and several common headers.

[Cancel](#) [Create resource](#)

NOTE: In current deployment resource path is "/health/" and resource name is "im"

c) After, create a POST method in that resource with following options :

[API Gateway](#) > [APIs](#) > [Resources - sagemaker \(fluqdn751l\)](#) > [Create method](#)

Create method

Method details

Method type

POST

Integration type

☐ Lambda function

Integrate your API with a Lambda function.



☐ HTTP

Integrate with an existing HTTP endpoint.



☐ Mock

Generate a response based on API Gateway mappings and transformations.



☒ AWS service

Integrate with an AWS Service.



☐ VPC link

Integrate with a resource that isn't accessible over the public internet.



AWS Region

us-east-1

AWS service

SageMaker Runtime

AWS subdomain

HTTP method

Action type

☐ Use action name

☒ Use path override

Path override - *optional*

endpoints/im/async-invocations

Execution role

arn:aws:iam::845842914165:role/service-role/lambda_apigateway

Credential cache

Do not add caller credentials to cache key

☒ Default timeout

“AWS service” as Integration type,

AWS Region is “us-east-1”,
 AWS service is “SageMaker Runtime” ,
 select “use path override” in Action type,
 in path override use: endpoint/<endpoint_name>/async-invocations,
 in Execution role use: arn:aws:iam::<account_id>:role/service-role/lambda_apigateway
 and save it.

NOTE: in current deployment, <endpoint_name> = im, this has to be same as the argument given while running python deploy.py –endpoint <endpoint_name> command previously

- d) Edit method request and go to HTTP request headers template add the following headers and save it.

Name	Required	Caching	
input	<input checked="" type="checkbox"/>	<input type="checkbox"/>	Remove
timeout	<input type="checkbox"/>	<input type="checkbox"/>	Remove

Add header

- e) Edit Integration request and go to URL request headers parameters template and add the following parameters and save it.

Name	Mapped from	Caching	
X-Amzn-SageMaker-Cust	method.request.body	<input type="checkbox"/>	Remove
X-Amzn-SageMaker-Inpu	method.request.header.ir	<input type="checkbox"/>	Remove
X-Amzn-SageMaker-Invo	method.request.header.ti	<input type="checkbox"/>	Remove

Add request header parameter

Name	Mapped from
X-Amzn-SageMaker-Custom-Attributes	method.request.body
X-Amzn-SageMaker-InvocationTimeoutSeconds	method.request.header.timeout

f) Deploy the API now.

API Gateway > APIs > Resources - sagemaker (fluqdn751l)

Resources

Create resource

- /
- /{model_name}
- POST
- /health
- /im (selected)
- POST

/health/im - POST - Method execution

Update documentation Delete

ARN: `arn:aws:execute-api:us-east-1:845842914165:fluqdn751l/*/POST/health/im` Resource ID: `8ogb0e`

Client → Method request → Integration request → AWS integration → Integration response → Method response → Client

g) Invoke the API by copying the URL

API Gateway > APIs > sagemaker (fluqdn751l) > Stages

Stages

Stage actions Create stage

- v1
- v2 (expanded)
- /health/im (selected)
- POST
- /{model_name}

Method overrides

Edit

By default, methods inherit stage-level settings. To customize settings for a method, configure method overrides.

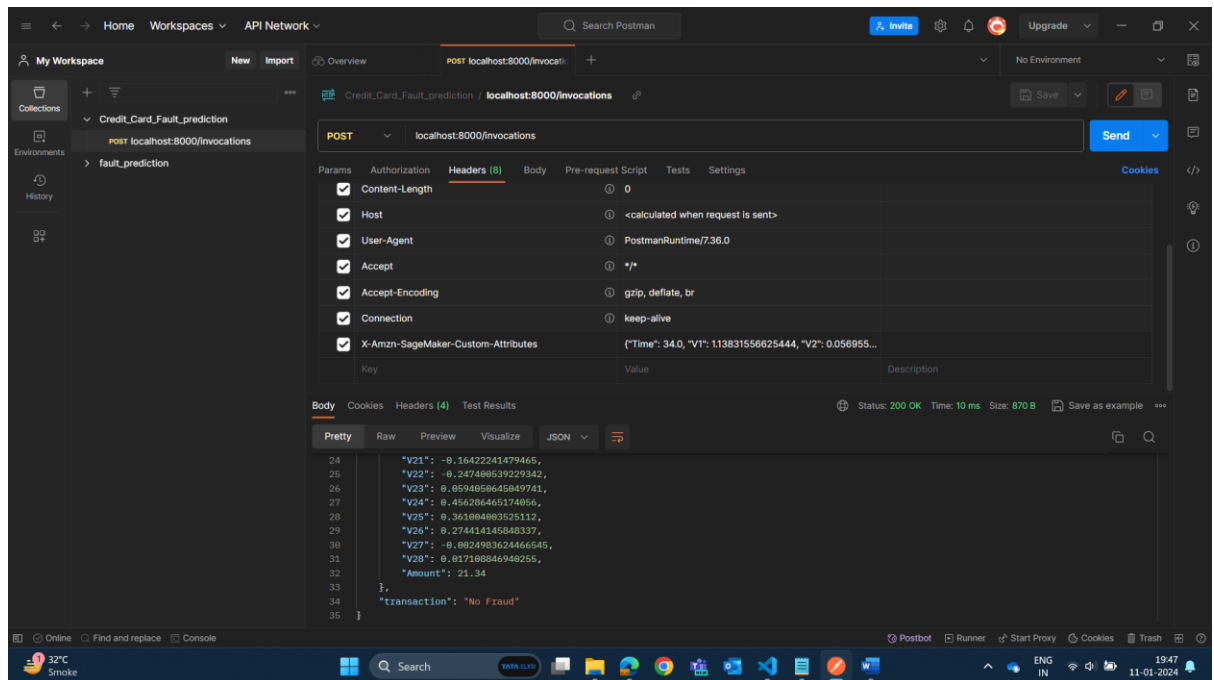
This method inherits its settings from the 'v2' stage.

Copied

`https://fluqdn751l.execute-api.us-east-1.amazonaws.com/v2/health/im`

8) Invoke the model using the created URL in API Gateway.

Local Deployment using Fast API and POSTMAN:

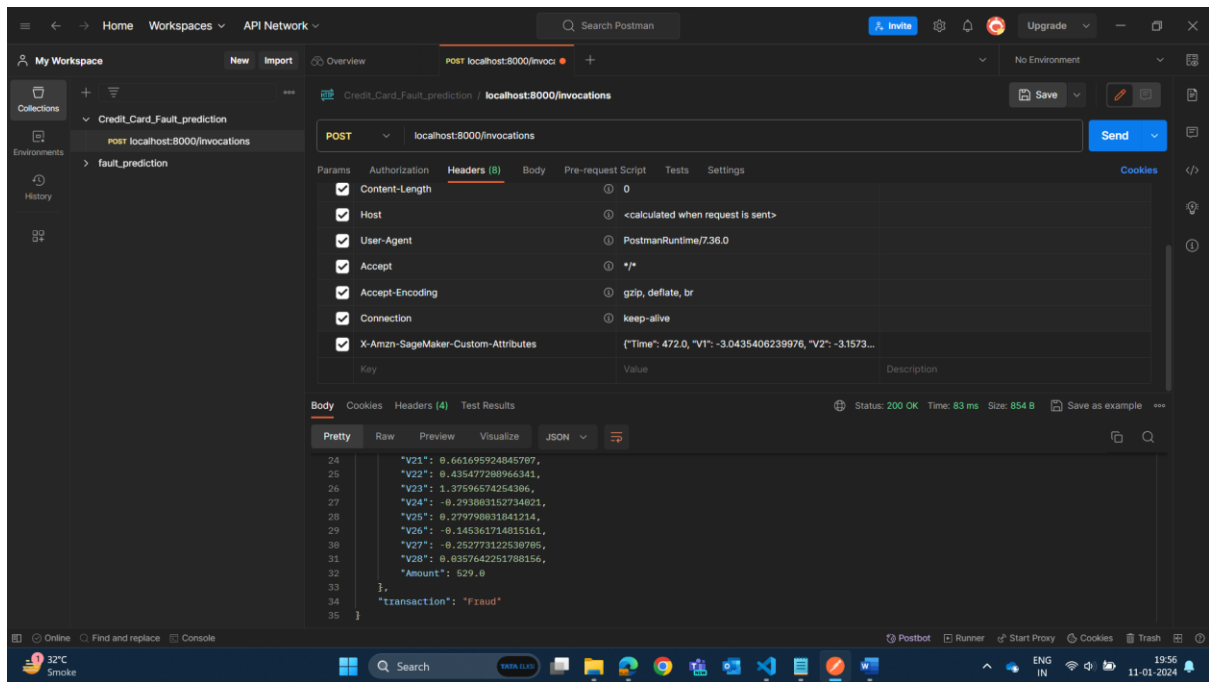


Logs showing the prediction which was invoked above: (Predicted – No Fault as No Fault)

```
INFO: 127.0.0.1:38720 - "POST /invocations HTTP/1.1" 200 OK
spec: {"Time": 34.0, "V1": 1.13831556625444, "V2": 0.056955969973862, "V3": 0.649418964599217, "V4": 0.873062048924213, "V5": -0.468466338715866, "V6": -0.410194552246249, "V7": -0.0138975543027732,
"V8": -0.0724396093157116, "V9": 0.306787909036459, "V10": -0.269952637390657, "V11": -0.0026918478230343, "V12": 1.12430421589913, "V13": 0.744206684168937, "V14": -0.188353421358499, "V15": -0.07566
196255573, "V16": -0.70919227171914, "V17": 0.381219102546287, "V18": -1.37288017404084, "V19": -0.273701149259017, "V20": -0.078354632808708, "V21": -0.16422241479465, "V22": -0.247480539229342, "V23":
0.0594050645049741, "V24": 0.456286465174056, "V25": 0.361004003525112, "V26": 0.274414145848337, "V27": -0.0024983624466545, "V28": 0.017108846940255, "Amount": 21.34}
type: <class 'dict'>
***** ENTERING fetch class *****
***** ENTERING predict class *****
/mnt/d/OneDrive - Tata Elxsi/[37184] Motor Cloud/cc/.venv/lib/python3.10/site-packages/sklearn/utils/validation.py:767: FutureWarning: is_sparse is deprecated and will be removed in a future version. C
heck 'isinstance(dtype, pd.SparseDtype)' instead.
if not hasattr(array, "sparse") and array.dtypes.apply(is_sparse).any():
/mnt/d/OneDrive - Tata Elxsi/[37184] Motor Cloud/cc/.venv/lib/python3.10/site-packages/sklearn/utils/validation.py:695: FutureWarning: is_sparse is deprecated and will be removed in a future version. C
heck 'isinstance(dtype, pd.SparseDtype)' instead.
if is_sparse(pd_dtype):
/mnt/d/OneDrive - Tata Elxsi/[37184] Motor Cloud/cc/.venv/lib/python3.10/site-packages/sklearn/utils/validation.py:614: FutureWarning: is_sparse is deprecated and will be removed in a future version. C
heck 'isinstance(dtype, pd.SparseDtype)' instead.
if is_sparse(pd_dtype) or not is_extension_array_dtype(pd_dtype):
INFO: 127.0.0.1:38720 - "POST /invocations HTTP/1.1" 200 OK
```

RESPONSE BODY of above request:

```
{
  "features": {
    "Time": 34.0,
    "V1": 1.13831556625444,
    "V2": 0.0569559699973862,
    "V3": 0.649418964599217,
    "V4": 0.873062040924213,
    "V5": -0.468466330715866,
    "V6": -0.410194552246249,
    "V7": -0.0138975543027732,
    "V8": -0.0724396093157116,
    "V9": 0.306787909036459,
    "V10": -0.269952637390657,
    "V11": -0.0026018478230343,
    "V12": 1.12430421589913,
    "V13": 0.744206684168937,
    "V14": -0.188353421358499,
    "V15": -0.07566196255573,
    "V16": -0.70919227171914,
    "V17": 0.381219102546287,
    "V18": -1.37208017404084,
    "V19": -0.273701149259017,
    "V20": -0.078354632808708,
    "V21": -0.16422241479465,
    "V22": -0.247400539229342,
    "V23": 0.0594050645049741,
    "V24": 0.456286465174056,
    "V25": 0.361004003525112,
    "V26": 0.274414145848337,
    "V27": -0.0024983624466545,
    "V28": 0.017108846940255,
    "Amount": 21.34
  },
  "transaction": "No Fraud"
}
```



Logs showing the prediction which was invoked above: (Predicted – Fault as Fault)

```
INFO: 127.0.0.1:38720 - "POST /invocations HTTP/1.1" 200 OK
specs: {'Time': 472.0, 'V1': -3.0435406239976, 'V2': -3.15730712090228, 'V3': 1.08846277997285, 'V4': 2.2886436183814, 'V5': 1.35980512966107, 'V6': -1.06482252298131, 'V7': 0.325574266158614, 'V8': -0.0677936531906277, 'V9': -0.270952836226548, 'V10': -0.838586564582682, 'V11': -0.414575448285725, 'V12': -0.503140859566824, 'V13': 0.676501544635863, 'V14': -1.69202893305906, 'V15': 2.00063483909015, 'V16': 0.666779695901966, 'V17': 0.599717413841732, 'V18': 1.72532100745514, 'V19': 0.283344830149495, 'V20': 2.10233879259444, 'V21': 0.661695924845707, 'V22': 0.435477208966341, 'V23': 1.37596574254306, 'V24': -0.293803152734021, 'V25': 0.279798031841214, 'V26': -0.145361714815161, 'V27': -0.252773122530705, 'V28': 0.8357642251788156, 'Amount': 529.0}
type: <class 'dict'>
***** ENTERING fetch class *****
***** ENTERING predict class *****
/mnt/d/OneDrive - Tata Elxsi/[37184] Motor Cloud/cc/.venv/lib/python3.10/site-packages/sklearn/utils/validation.py:767: FutureWarning: is_sparse is deprecated and will be removed in a future version. C
heck 'isinstance(dtype, pd.SparseDtype)' instead.
if not hasattr(array, "sparse") and array.dtypes.apply(is_sparse).any():
/mnt/d/OneDrive - Tata Elxsi/[37184] Motor Cloud/cc/.venv/lib/python3.10/site-packages/sklearn/utils/validation.py:695: FutureWarning: is_sparse is deprecated and will be removed in a future version. C
heck 'isinstance(dtype, pd.SparseDtype)' instead.
if is_sparse(pd_dtype):
/mnt/d/OneDrive - Tata Elxsi/[37184] Motor Cloud/cc/.venv/lib/python3.10/site-packages/sklearn/utils/validation.py:614: FutureWarning: is_sparse is deprecated and will be removed in a future version. C
heck 'isinstance(dtype, pd.SparseDtype)' instead.
if is_sparse(pd_dtype) or not is_extension_array_dtype(pd_dtype):
INFO: 127.0.0.1:40644 - "POST /invocations HTTP/1.1" 200 OK
```

RESPONSE BODY of above request:

```
{
  "features": {
    "Time": 472.0,
    "V1": -3.0435406239976,
    "V2": -3.15730712090228,
    "V3": 1.08846277997285,
    "V4": 2.2886436183814,
    "V5": 1.35980512966107,
    "V6": -1.06482252298131,
    "V7": 0.325574266158614,
    "V8": -0.0677936531906277,
    "V9": -0.270952836226548,
    "V10": -0.838586564582682,
    "V11": -0.414575448285725,
    "V12": -0.503140859566824,
    "V13": 0.676501544635863,
    "V14": -1.69202893305906,
    "V15": 2.00063483909015,
    "V16": 0.666779695901966,
    "V17": 0.599717413841732,
    "V18": 1.72532100745514,
    "V19": 0.283344830149495,
    "V20": 2.10233879259444,
    "V21": 0.661695924845707,
    "V22": 0.435477208966341,
    "V23": 1.37596574254306,
    "V24": -0.293803152734021,
    "V25": 0.279798031841214,
    "V26": -0.145361714815161,
    "V27": -0.252773122530705,
    "V28": 0.0357642251788156,
    "Amount": 529.0
  },
  "transaction": "Fraud"
}
```


Future Roadmap

3.1) Ensemble model prediction

To go through the diverse algorithms, an ensemble approach is adopted which will include models like Adaboost, XGboost, Catboost, LGBM. Combining diverse models can capture different patterns in the data. And later, stacking will be done, stacking often leads to better predictive performance compared to individual models.

3.2) Aggregation of Analytical and Machine Learning Approach

Implement a voting mechanism to aggregate predictions.

3.2.1) Voting based aggregation:

Predicted class is the one that receives most votes from the models.

3.2.2) Weightage based aggregation:

Assign different weights of the predictions of each model based on their performance or reliability.